

MACHINE LEARNING

William Lowe

Hertie School

27th September 2020

MACHINE LEARNING: RATHER QUICKLY

Plan

- Types of machine learning
- Going further than regression
- More flexibility with polynomials
- Overfitting and regularization
- Bias variance decomposition
- The good part about bias
- Example ML models: Ridge, lasso, trees, neural networks
- Back to causal inference in the style of Frisch, Waugh, and Lovell
- Double/debiased ML, the very idea

MACHINE LEARNING: RATHER QUICKLY

Inference:

- Supervised: Learn $P(Y \mid X, Z \dots)$, or often just its expected value
- Unsupervised: Learn $P(X, Z)$

MACHINE LEARNING: RATHER QUICKLY

Inference:

- Supervised: Learn $P(Y \mid X, Z \dots)$, or often just its expected value
- Unsupervised: Learn $P(X, Z)$

Action (embeds an inference problem)

- Reinforcement: Learn a policy $P(\text{Action} \mid \text{State})$ such that the expected future discounted *reward* for the policy's actions is maximized

MACHINE LEARNING: RATHER QUICKLY

Inference:

- Supervised: Learn $P(Y \mid X, Z \dots)$, or often just its expected value
- Unsupervised: Learn $P(X, Z)$

Action (embeds an inference problem)

- Reinforcement: Learn a policy $P(\text{Action} \mid \text{State})$ such that the expected future discounted *reward* for the policy's actions is maximized

We'll be interested in supervised, traditionally separated into

- Regression: usually implicitly assumes symmetric constant ϵ (or doesn't have an opinion...)
- Classification: ambiguous between *choosing* one of K classes and estimating $P(Y = k \mid X, Z \dots)$

In any case, both go for $E[Y \mid X, Z \dots]$

FLEXIBILITY

You could, if you like, think of linear regression and logistic regression in each of these categories

→ It's illuminating to do so (see the first few chapters of Bishop, 2006)

So what's the difference?

→ More flexible forms for $E[Y \mid X, Z \dots]$

→ Higher dimensional predictors, i.e. lots more $X, Z \dots$

Many ML regression models will embed a more familiar model, e.g. neural networks

Others will start from scratch and build $E[Y \mid X, Z \dots]$ differently, e.g. classification trees

INDIFFERENCE

As an engineering tool, ML models will seldom care about what X , Z etc. actually are, or distinguishing one parameter among the others

Indeed most are *non-parametric*

- Reminder: ‘non-parametric’ does not mean ‘does not have parameters, it means ‘has so many parameters that I do not care to know them by name’

Unsurprisingly, this part of ML came late to causal inference

EXTENDING REGRESSION

What happens when there are more variables than cases?

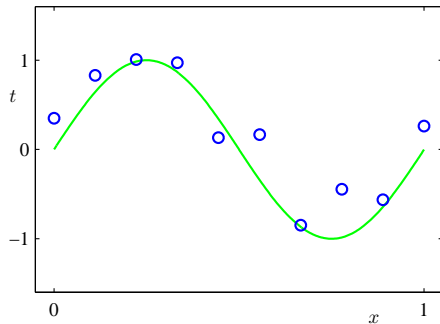
- Regular regression breaks

What happens when you add all the squares and cubes and interactions as predictors?

- Standard errors explode; same amount of data, but more parameters to learn from it.
- Generalization to new data gets worse; now we can fit everything better, we fit noise better

These are the same problem in different degrees

ADDING POLYNOMIALS



For consistency with Bishop ch. 1 let's call

- the outcome t_n
- the regression coefficients $w_j \in \mathbf{w}$ ('weights')
- our estimate of the expected value of t_n , $\hat{t}_n = y(x, \mathbf{w})$

ADDING FLEXIBILITY

Consider polynomial models of t . We'll fit / make predictions like this:

$$\begin{aligned}y(x, \mathbf{w}) &= w_0 \\&= w_0 + w_1 x \\&= w_0 + w_1 x + w_2 x^2 \\&= w_0 + w_1 x + w_2 x^2 + \cdots + w_M x^M \\&= \sum_j^M w_j x^j\end{aligned}$$

ADDING FLEXIBILITY

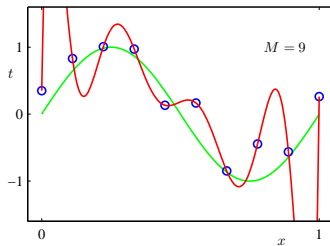
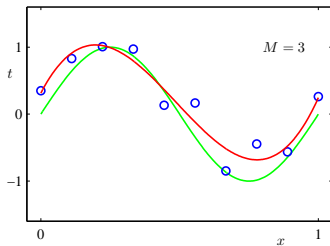
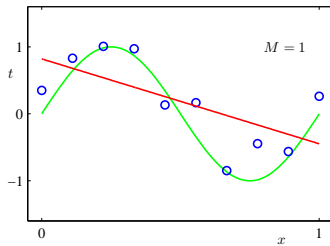
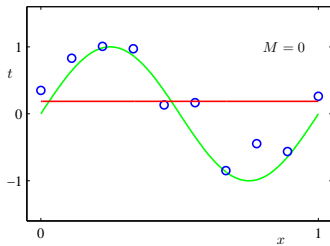
Consider polynomial models of t . We'll fit / make predictions like this:

$$\begin{aligned}y(x, \mathbf{w}) &= w_0 \\&= w_0 + w_1 x \\&= w_0 + w_1 x + w_2 x^2 \\&= w_0 + w_1 x + w_2 x^2 + \cdots + w_M x^M \\&= \sum_j^M w_j x^j\end{aligned}$$

The *flexibility* of this model is driven by M , which we can think of as determining the model class

→ Roughly: the set of functions that can be represented

ADDING FLEXIBILITY

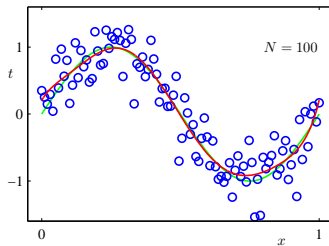
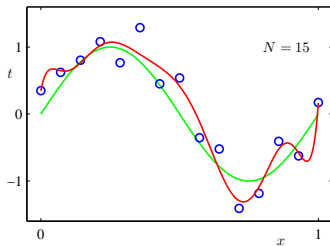


ADDING FLEXIBILITY

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

OVERFITTING

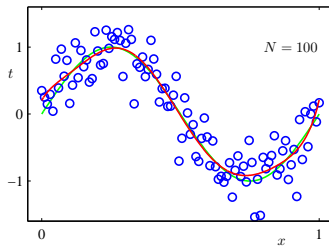
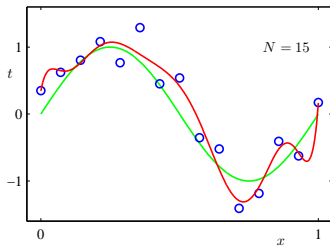
Things are not so bad when there is more data (here $M=9$)



But there isn't always going to be more data...

OVERFITTING

Things are not so bad when there is more data (here $M=9$)

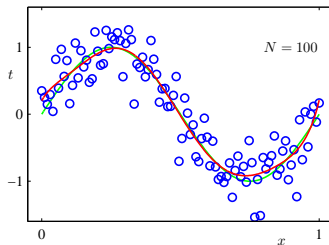
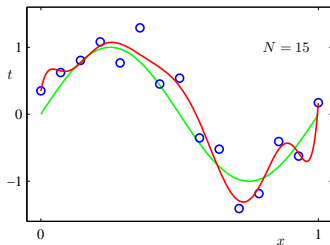


But there isn't always going to be more data...

However, we can keep all M , i.e. maintain the flexibility in the model class, *if* we can constrain the size of the weights

→ This calls for a *hyperparameter*, a parameter that controls other parameters

OVERFITTING



When there's lots of persuasive data:

→ override the hyperparameter and make use of the model flexibility

When there isn't,

→ keep weights small, and therefore the function smooth

REGULARIZATION BY HYPERPARAMETER

Here, we're fitting the model (maximizing the likelihood) using OLS, which *minimises* the sum of squared errors

$$E_{\text{OLS}} = \frac{1}{2} \sum_n^N (y(x_n, \mathbf{w}) - t_n)^2$$

Note: minimising error rather than maximizing the likelihood is the way ML people think about things (hence, no minus sign)¹

¹The 1/2 is there to hint that this is the log likelihood for a Normal distribution (with constant error variance, so it doesn't matter to E)

REGULARIZATION BY HYPERPARAMETER

Here, we're fitting the model (maximizing the likelihood) using OLS, which *minimises* the sum of squared errors

$$E_{\text{OLS}} = \frac{1}{2} \sum_n^N (y(x_n, \mathbf{w}) - t_n)^2$$

Note: minimising error rather than maximizing the likelihood is the way ML people think about things (hence, no minus sign)¹

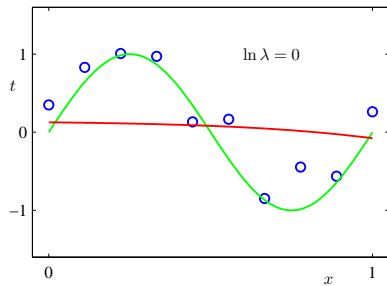
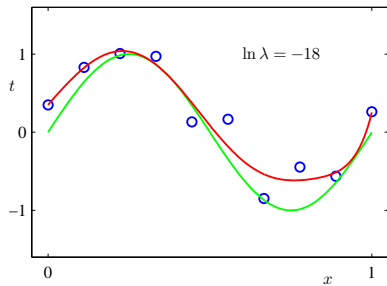
Let's keep that plan, but add an extra term to control the weights

$$E_\lambda = \frac{1}{2} \sum_n^N (y(x_n, \mathbf{w}) - t_n)^2 + \frac{\lambda}{2} \sum_m^M w_m^2$$

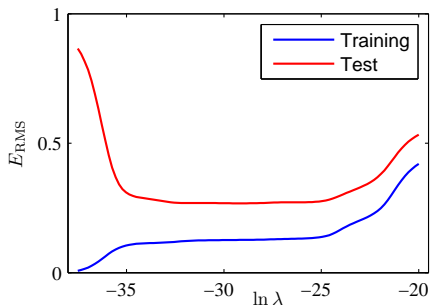
and a hyperparameter λ to say how seriously we should take it as an error component

¹The 1/2 is there to hint that this is the log likelihood for a Normal distribution (with constant error variance, so it doesn't matter to E)

CONSEQUENCES



THE SWEET SPOT



- The left extreme is $\lambda = 0$ (no regularization)
- The right extreme is all zero weights (predict of 0 for every point)
- With fixed data, decreasing λ allows more of the model class's inherent flexibility to show

CHOOSING HYPERPARAMETER VALUES

We can't fit λ by minimising the sum of squares

→ That would just set it to zero (why?)

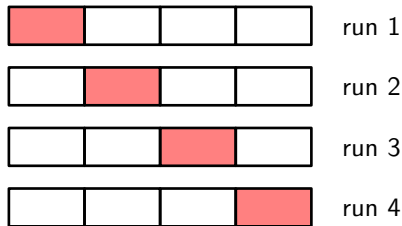
CHOOSING HYPERPARAMETER VALUES

We can't fit λ by minimising the sum of squares

→ That would just set it to zero (why?)

One reliable option is crossvalidation (CV)

- Make a grid of hyperparameter values
- Randomly divide the data set into 4 (or some other value > 1)
- For each hyperparameter value, train a model on white and test on red
- Choose the hyperparameter value that minimizes the average error on reds



GENERAL THEORY: BIAS-VARIANCE

One important question we can ask is about the expected value of E , averaged over *all possible data sets* coming from the same mechanisms

First, let's give $E[t \mid x]$ (the *real* regression function) a name

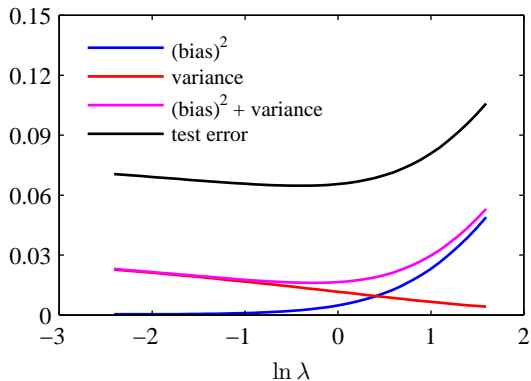
$$h(x) = E[t \mid x]$$

and define E_D as an average over all possible data sets

Then (Bishop, sec. 1.5.5 and 3.2 for a derivation) we can decompose the expected error into

$$\begin{aligned} & \text{bias}^2 \quad \int E_D[(y(x, \mathbf{w}) - h(x))^2] p(x) dx \\ & \text{variance} \quad \int E_D[(y(x, \mathbf{w}) - E_D[y(x, \mathbf{w})])^2] p(x) dx \\ & \text{noise} \quad \int (h(x) - t)^2 p(x, t) dx dt \end{aligned}$$

BIAS AND VARIANCE



Error is unavoidable but bias and variance trade off

MODEL BIAS

Informally, the *bias* of a model class is the set of functions that a model most naturally learns

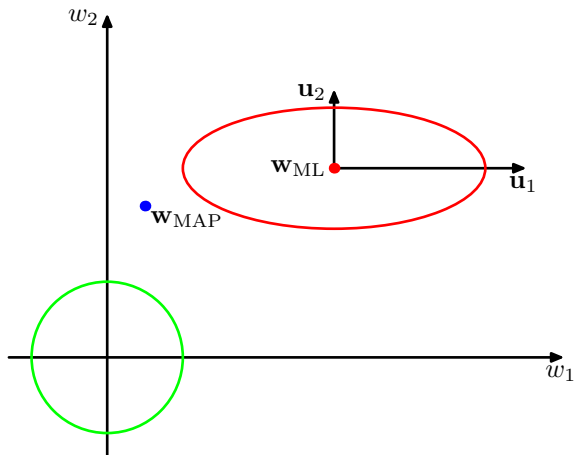
- Linear models ($M=1$ above): can learn straight lines
- Quadratic models ($M=2$ above): *can* learn straight lines but also smooth curves
- etc.

We can get different sorts of bias by changing the whole model class

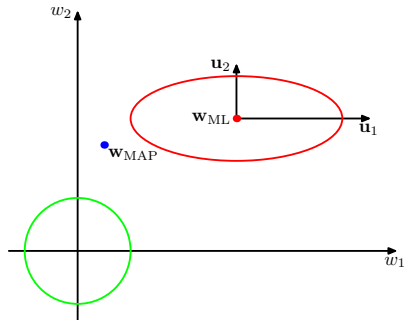
- we'll see an example of later with trees

Regularization also offers us some interesting and different forms

MODEL BIAS

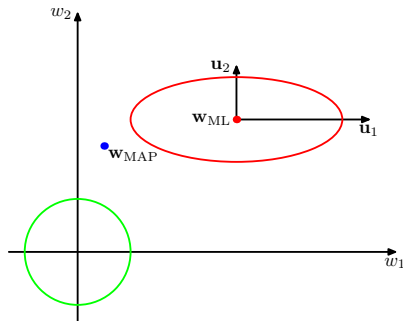


MODEL BIAS



- \mathbf{w}_{ML} minimizes $E_{\lambda=0} = E_{\text{OLS}}$
- The origin minimizes $E_{\lambda=\infty}$
- \mathbf{w}_{MAP} minimizes E_{λ} when we set λ sensibly to balance the two parts of the error function

MODEL BIAS



- This bias *shrinks* all the weights, some more than others
- It is sometimes helpful to define an *effective* number of parameters, which is less than M , and possibly fractional

A DIFFERENT MODEL BIAS

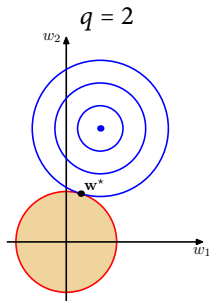
If we change the regularization term just a little (note the q)

$$E_{\lambda} = \frac{1}{2} \sum_n^N (y(x_n, \mathbf{w}) - t_n)^2 + \frac{\lambda}{2} \sum_m^M |w_m|^q$$

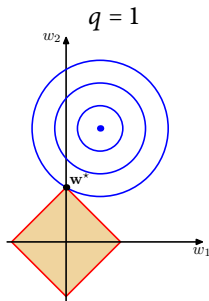
A DIFFERENT MODEL BIAS

If we change the regularization term just a little (note the q)

$$E_{\lambda} = \frac{1}{2} \sum_n^N (y(x_n, \mathbf{w}) - t_n)^2 + \frac{\lambda}{2} \sum_m^M |w_m|^q$$



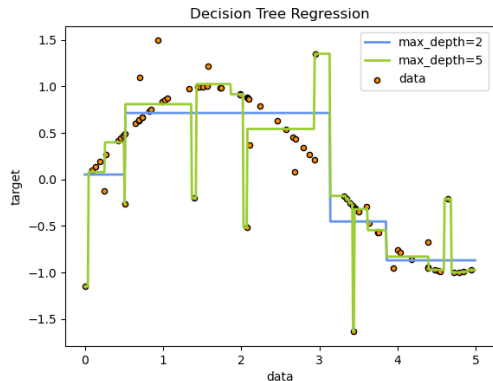
(L2 regularization a.k.a. ‘ridge regression’)



(L1 regularization, a.k.a. ‘the lasso’)

TOTALLY DIFFERENT BIAS

Alternatively, we can change the model class altogether, e.g. regression trees (from scikit-learn documentation)



Blue tree (2 levels)

```
if x > 3.2 then
    if x > 3.9 then -0.9 else -0.5
else
    if x > 0.5 then 0.8 else 0.1
```

The green tree allows up to 5 levels, and overfits

REGRESSION TREES

For regression trees, one hyperparameter is the depth of the tree

- so constraining that adds bias and reduces variance

In general we can also prevent overfitting by bagging (Breiman, 1996)

- bootstrapping the dataset
- Fitting trees to each bootstrap sample
- Averaging the resulting predictions

or variations on that theme (Cutler et al., 2012, e.g. Random forests)

Like cross-validation, this removes variance but does not much affect bias

BIAS AS A GOOD THING

Clearly regularization generates bias. Seems like a bad thing...

But it's necessary

- The *No Free Lunch theorem* (Wolpert, 1996) says that averaged over all possible problems, no learning algorithm is better than any other
- Happily we don't deal with all possible problems, so we can and should choose a model bias to fit the problem

BIAS AS A GOOD THING

Clearly regularization generates bias. Seems like a bad thing...

But it's necessary

- The *No Free Lunch theorem* (Wolpert, 1996) says that averaged over all possible problems, no learning algorithm is better than any other
- Happily we don't deal with all possible problems, so we can and should choose a model bias to fit the problem

And helpful

- It's how we get less variance

BIAS AS A GOOD THING

Clearly regularization generates bias. Seems like a bad thing...

But it's necessary

- The *No Free Lunch theorem* (Wolpert, 1996) says that averaged over all possible problems, no learning algorithm is better than any other
- Happily we don't deal with all possible problems, so we can and should choose a model bias to fit the problem

And helpful

- It's how we get less variance

And annoying

- Slows convergence

This is better than the alternative, which is not being consistent and not knowing it

BIAS AND VARIANCE IN ML

ML insight:

- It's better to working with a universal function approximator, and figure out how to regularize it, than to work with a model that can't represent much of anything and hope
- Most of the ML methods we'll work with are universal approximators
- Linear regression... definitely not.

BIAS AND VARIANCE IN ML

ML insight:

- It's better to working with a universal function approximator, and figure out how to regularize it, than to work with a model that can't represent much of anything and hope
- Most of the ML methods we'll work with are universal approximators
- Linear regression... definitely not.

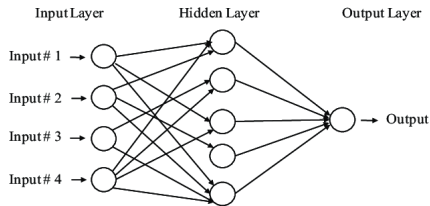
But wait, how did all that regularization business turn $y(x, \mathbf{w})$ into a universal approxiator?

- It didn't. We just didn't say much about $y(x, \mathbf{w})$ and drew it like it was in a linear regression context

In real applications, $y(x, \mathbf{w})$ is not even polynomial regression

- It's *kernel regression*, or *basis function regression*, *neural network*, *random forest*, etc.

NEURAL NETWORKS



A multilayer perceptron (MLP) with 1 hidden layer of J 'units' for D -dimensional input data x is

$$y(x, \mathbf{w}) = \sum_j^J w_j \phi_j(x, \mathbf{w}^{(j)})$$

where ϕ_j is some nonlinear function of the input data, e.g.

$$\phi_j(x, \mathbf{w}^{(j)}) = 1/(1 + \exp(-\sum_d w_d^{(j)} x_d))$$

That's a universal approximator (Hornik et al., 1989) that needs serious regularization

ML MODELS

We've some really flexible models with interesting different types of bias (smooth, piecewise linear) and styles of regularization (L1, L2, depth constraints)

Let's go right back to the beginning

→ Good old multiple linear regression

ML MODELS

We've some really flexible models with interesting different types of bias (smooth, piecewise linear) and styles of regularization (L1, L2, depth constraints)

Let's go right back to the beginning

→ Good old multiple linear regression

EXTENDING REGRESSION

Frisch and Waugh, back in 1933 in *Econometrica* showed, and Lowell generalized in 1963 the following useful fact about regression (Lovell, 2008, has a short accessible proof).

Consider three models

$$Y = \beta_0 + X\beta_X + Z_1\beta_{Z1} \dots Z_K\beta_{ZK} + \epsilon \quad (\text{Big Model})$$

$$X = \beta_0^X + Z_1\beta_{Z1}^X \dots Z_K\beta_{ZK}^X + \epsilon^X \quad (\text{X Model})$$

$$Y = \beta_0^Y + Z_1\beta_{Z1}^Y \dots Z_K\beta_{ZK}^Y + \epsilon^Y \quad (\text{Y Model})$$

and also this one, made out of residuals from the Y and X models

$$(Y - \hat{Y}) = \beta_0^{\text{FWL}} + (X - \hat{X})\beta^{\text{FWL}} + \epsilon^{\text{FWL}}$$

then

$$\beta_X = \beta^{\text{FWL}}$$

EXTENDING REGRESSION

What if K were really big and we had no real idea about all those Z s?

- Lots of confounders
- Unknown possibly non-linear relationships
- Target effect is β_X , which we'll assume here *is* linear

EXTENDING REGRESSION

What if K were really big and we had no real idea about all those Z s?

- Lots of confounders
- Unknown possibly non-linear relationships
- Target effect is β_X , which we'll assume here *is* linear

Double/debiased Machine Learning (Chernozhukov et al., 2018)

Intuition: Do FWL but with *fancier* X and Y models

$$X = m(Z_1, \dots, Z_K) + \epsilon^X \quad (\text{Fancy X Model})$$

$$Y = g(Z_1, \dots, Z_K) + \epsilon^Y \quad (\text{Fancy Y Model})$$

DOUBLE ML

Chernozhukov et al. show that

- Just learning a fancier big model is a bad idea: the effect of X gets lost, the fancy model might throw it away, bias, etc.

However using fancy m and fancy g comes with problems:

- Overfitting, due to flexibility of the model class
- Bias, due regularization to combat overfitting
- Slow convergence. we're used to $n^{1/2}$, but fancy models tend to go $n^{1/4}$

They use a mixture of

- cross-fitting: like cross validation but for β_X estimation
- cunning orthogonal score functions

To get fancy models that converge (mostly) as if they were simple ones. Cool

REFERENCES

- Bishop, C. M. (2006). 'Pattern recognition and machine learning'. Springer.
- Breiman, L. (1996). 'Bagging predictors'. *Machine Learning*, 24(2), 123–140.
- Chernozhukov, V., Chetverikov, D., Demirer, M., Duflo, E., Hansen, C., Newey, W. & Robins, J. (2018). 'Double/debiased machine learning for treatment and structural parameters'. *The Econometrics Journal*, 21(1), C1–C68.
- Cutler, A., Cutler, D. R. & Stevens, J. R. (2012). Random forests. In C. Zhang & Y. Ma (Eds.), *Ensemble machine learning* (pp. 157–175). Springer US.
- Hornik, K., Stinchcombe, M. & White, H. (1989). 'Multilayer feedforward networks are universal approximators'. *Neural Networks*, 2, 359–366.
- Lovell, M. C. (2008). 'A simple proof of the fwl theorem'. *The Journal of Economic Education*, 39(1), 88–91
_eprint: <https://doi.org/10.3200/JECE.39.1.88-91>.

REFERENCES

Wolpert, D. H. (1996). ‘The lack of a priori distinctions between learning algorithms’. *Neural Computation*, 8, 1341–1390.